

# Non-parametric Graph Convolution for Re-ranking in Recommendation Systems

Zhongyu Ouyang\*

zhongyu.ouyang.gr@dartmouth.edu  
Dartmouth College  
Hanover, New Hampshire, USA

Soroush Vosoughi‡

soroush.vosoughi@dartmouth.edu  
Dartmouth College  
Hanover, New Hampshire, USA

Mingxuan Ju†\*

mju@snap.com  
Snap Inc.

Bellevue, Washington, USA  
University of Notre Dame  
Notre Dame, Indiana, USA

Yanfang Ye‡

yee7@nd.edu  
University of Notre Dame  
Notre Dame, Indiana, USA

## Abstract

Graph knowledge has been proven effective in enhancing item rankings in recommender systems (RecSys), particularly during the retrieval stage. However, its application in the ranking stage, especially when richer contextual information in user-item interactions is available, remains underexplored. A major challenge lies in the substantial computational cost associated with repeatedly retrieving neighborhood information from billions of items stored in distributed systems. This resource-intensive requirement makes it difficult to scale graph-based methods in practical RecSys. To bridge this gap, we first demonstrate that incorporating graphs in the ranking stage improves ranking qualities. Notably, while the improvement is evident, we show that the substantial computational overheads entailed by graphs are prohibitively expensive for real-world recommendations. In light of this, we propose a non-parametric strategy that utilizes graph convolution for re-ranking only during test time. Our strategy circumvents the notorious computational overheads from graph convolution during training, and utilizes structural knowledge hidden in graphs on-the-fly during testing. It can be used as a plug-and-play module and easily employed to enhance the ranking ability of various ranking layers of a real-world RecSys with significantly reduced computational overhead. Through comprehensive experiments across four benchmark datasets with varying levels of sparsity, we demonstrate that our strategy yields noticeable improvements (i.e., **8.1%** on average) during testing time with little to no additional computational overheads (i.e., **0.5%** on average). Code: [https://github.com/zyouyang/RecSys2025\\_NonParamGC.git](https://github.com/zyouyang/RecSys2025_NonParamGC.git)

\*Equal contributions.

†This work does not relate to his role at Snap Inc.

‡Corresponding author.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

RecSys '25, Prague, Czech Republic

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.  
ACM ISBN 979-8-4007-1364-4/2025/09  
<https://doi.org/10.1145/3705328.3748058>

## Keywords

Recommender Systems, Re-ranking, Efficient Machine Learning Systems, Test-Time Augmentation

### ACM Reference Format:

Zhongyu Ouyang, Mingxuan Ju, Soroush Vosoughi, and Yanfang Ye. 2025. Non-parametric Graph Convolution for Re-ranking in Recommendation Systems. In *Proceedings of the Nineteenth ACM Conference on Recommender Systems (RecSys '25), September 22–26, 2025, Prague, Czech Republic*. ACM, New York, NY, USA, 10 pages. <https://doi.org/10.1145/3705328.3748058>

## 1 Introduction

Recommender systems (RecSys) are central to user experience in diverse online applications, ranging from product recommendation on e-commerce platforms [26, 31, 39] and personalized advertising [6, 15, 25, 35], to friend recommendation on social media [5, 13, 14, 23]. Typical recommender systems follow a two-stage pipeline: *retrieval* and *ranking*. The *retrieval* stage filters a massive item pool, often billions of items, down to a manageable subset using lightweight models or approximate matching. This stage commonly leverages collaborative filtering (CF), which utilizes historical user-item interactions to recommend items based on similar behavioral patterns [16, 19, 30, 37]. Classical CF approaches typically employ matrix factorization [30, 37], representing users and items with learned embeddings to reconstruct historical interactions.

In contrast, the *ranking* stage uses more sophisticated models incorporating rich contextual features such as user demographics, item attributes, and interaction details like transaction timestamps. Context-aware ranking models better capture user preference variability across different contexts. To balance performance and efficiency, the ranking problem is typically formulated as a click-through rate (CTR) prediction task [8, 11, 38], simplified as predicting binary interaction outcomes. The predicted probabilities serve as scores for ranking candidate items.

Recently, graph neural networks (GNNs) have emerged as powerful methods for modeling relational data, inspiring extensive research on user-item interaction bipartite graphs in RecSys [12, 27, 40, 45]. GNN-based methods like NGCF [40], LightGCN [12], and SimGCL [45] have shown significant improvements, notably in the retrieval stage, by leveraging high-order connectivity within

interaction graphs, achieving performance gains up to approximately 40% compared to traditional methods [40]. Nevertheless, the integration of GNNs into the *ranking* stage remains underexplored, primarily due to the computational complexity of repeatedly querying graph neighborhoods, a bottleneck magnified by the reliance on detailed contextual data.

A recent attempt to incorporate graphs at the ranking stage, the Graph Convolution Machine (GCM), directly applies message passing over user-item bipartite graphs [42], enriching nodes and edges with contextual features. Despite improved ranking quality, GCM faces significant hurdles in practical deployment:

(i) GCM is **computationally prohibitive to train** at an industrial scale, where billions of users and items reside on distributed infrastructures. The message-passing operation necessitates querying neighborhood embeddings repeatedly, incurring quadratic computational complexity relative to batch sizes. This overhead is further exacerbated in distributed environments due to bandwidth constraints.

(ii) Integrating GCM with existing context-aware RecSys pipelines **requires extensive engineering effort**, as industrial systems predominantly utilize deep models like DCN [38] designed for tabular, independently and identically distributed (i.i.d.) data, contrasting sharply with non-i.i.d. graph structures. Therefore, it is critical to develop a practical method to efficiently leverage graph knowledge within existing industrial frameworks.

Consequently, this work addresses the following core research question:

### How can we effectively yet efficiently integrate graph knowledge into the *ranking* stage of existing RecSys?

To address this gap, we first empirically validate the potential benefits of combining contextual and structural information. Specifically, we perform a proof-of-concept experiment substituting traditional embedding tables in a widely used industrial ranking model with a graph encoder performing message passing. Although this approach significantly boosts performance across four benchmark datasets, it incurs substantial computational costs and requires infrastructural modifications, limiting practical applicability.

Motivated by these limitations, we propose a test-time augmentation strategy that incorporates graph structure only at inference, bypassing expensive training-phase computations. Our method seamlessly applies to any pre-trained ranking model with minimal additional overhead, involving four simple steps: (i) constructing graph-based similarity matrices between users and items; (ii) retrieving relevant candidate users and items based on these similarities; (iii) forming augmented user-item pairs to obtain predicted interaction probabilities from the target model; and (iv) aggregating these predictions through weighted fusion to generate a final ranking score. Extensive experimentation across multiple benchmark datasets demonstrates the efficacy and efficiency of our approach, achieving significant ranking improvements with negligible additional computational costs.

## 2 Non-parametric Graph Convolution for Re-ranking

Without loss of generality, we adopt a CTR model as our ranking model and first illustrate the standard paradigm in industrial RecSys, which does not incorporate graph information. Subsequently, we present a graph-based encoder designed to integrate structural user-item relationships. Although this encoder notably improves ranking performance, it introduces substantial computational overhead—up to approximately 1000% increase during training—which severely limits its practicality. To resolve this challenge, we propose a test-time graph augmentation strategy that leverages graph-based insights exclusively during inference, significantly enhancing ranking quality without incurring additional training costs. Figure 1 provides an overview of our proposed approach.

### 2.1 Industrial Ranking Paradigms

Formally, consider a user  $i$  and item  $j$  represented by their IDs  $x_i$  and  $x_j$ , respectively, along with contextual features associated with their interaction, denoted as  $c_{ij} \in \mathbb{R}^{d^c}$ , where  $d^c$  represents the contextual feature dimension. The encoder for user/item IDs is defined as  $f(\cdot) : \mathbb{R} \rightarrow \mathbb{R}^d$ , mapping IDs to latent embeddings of dimension  $d$ , and the contextual feature encoder as  $h(\cdot) : \mathbb{R}^{d^c} \rightarrow \mathbb{R}^{d'}$ , where  $d'$  is the dimensionality of the contextual embeddings.

In typical CTR models, ranking is formulated as a binary classification problem, aiming to predict the interaction probability between a given user-item pair within the range [0,1]. This process is illustrated in Figure 1(c). Specifically, the input vector  $\mathbf{z}_{ij} \in \mathbb{R}^{2d+d'}$  to a CTR model is constructed as:

$$\mathbf{z}_i = f(x_i), \quad \mathbf{z}_j = f(x_j), \quad \mathbf{z}_{ij}^c = h(c_{ij}), \quad (1)$$

$$\mathbf{z}_{ij} = [\mathbf{z}_i \parallel \mathbf{z}_j \parallel \mathbf{z}_{ij}^c], \quad (2)$$

where  $\parallel$  denotes concatenation. This embedding strategy is depicted in Figure 1(a,c).

The conventional ID encoder  $f(\cdot)$  used in ranking stages is typically a lookup embedding table represented as  $E \in \mathbb{R}^{(|\mathcal{U}|+|\mathcal{I}|) \times d}$ , with  $\mathcal{U}$  and  $\mathcal{I}$  denoting the sets of users and items, respectively. With the encoded embeddings, the CTR model applies a rating function  $r(\cdot) : \mathbb{R}^{2d+d'} \rightarrow \mathbb{R}$  to estimate the interaction likelihood, defined as  $p_{ij} = r(\mathbf{z}_{ij})$ .

To illustrate how contextual information is incorporated into ranking models, we highlight DCN [39], a widely used CTR model in industrial scenarios. DCN consists of two primary components: a cross network and a deep network. The cross network explicitly models feature interactions via cross layers:

$$\mathbf{z}^{(l+1)} = \mathbf{z}^{(0)} \mathbf{z}^{(l)\top} \mathbf{w}^{(l)} + \mathbf{b}^{(l)} + \mathbf{z}^{(l)}, \quad (3)$$

where  $\mathbf{z}^{(l)}, \mathbf{z}^{(l+1)} \in \mathbb{R}^d$  are the input and output vectors for the  $l$ -th cross layer, and  $\mathbf{w}^{(l)}, \mathbf{b}^{(l)} \in \mathbb{R}^d$  represent learnable parameters.

Conversely, the deep network captures complex, nonlinear interactions through fully connected layers:

$$\mathbf{h}^{(l+1)} = \text{ReLU}(\mathbf{W}^{(l+1)} \mathbf{h}^{(l)} + \mathbf{b}^{(l)}), \quad (4)$$

where  $\mathbf{h}^{(l)}, \mathbf{h}^{(l+1)}$  denote the input and output of the  $l$ -th hidden layer, and  $\mathbf{W}^{(l)}, \mathbf{b}^{(l)}$  are trainable parameters.




Figure 1: (a) Table encoder for user/item ID embeddings; (b) Graph-enhanced user/item ID encoder; (c) Ranking score calculation; (d) Proposed test-time graph augmentation strategy.

Given a user-item pair, the initial inputs are set as  $\mathbf{z}^{(0)} = \mathbf{h}^{(0)} = \mathbf{z}_{ij}$ . The outputs from the cross and deep networks are concatenated and passed through a two-class logits layer to yield the final prediction:

$$p_{ij} = \sigma([\mathbf{z}^{(L_1)} \parallel \mathbf{h}^{(L_2)}] \mathbf{w}_{\text{logits}}), \quad (5)$$

where  $\sigma(x) = 1/(1+e^{-x})$  and  $\mathbf{w}_{\text{logits}}$  are weights in the logits layer. The DCN model optimizes the binary cross-entropy loss:

$$\mathcal{L} = -\frac{1}{N} \sum_{\{i,j\} \in T_r} y_{ij} \log(p_{ij}) + (1-y_{ij}) \log(1-p_{ij}) + \lambda \sum_l \|\mathbf{w}^{(l)}\|^2, \quad (6)$$

with  $T_r$  denoting the training set of interaction pairs,  $y_{ij}$  being the binary labels, and  $\lambda$  the regularization coefficient.

## 2.2 A Naive Graph-based Ranking Framework

In this section, we first introduce the definition of the user-item interaction bipartite graph that depicts rich topological relationships such as co-purchase and shared interests. Then, we present how graph knowledge can be naively incorporated into existing ranking methods in their training processes.

Formally, let the interaction matrix be  $M \in \{0, 1\}^{|\mathcal{U}| \times |\mathcal{I}|}$ , where  $m_{ij} = 1$  represents an observed positive interaction between user  $i$  and item  $j$ , and  $m_{ij} = 0$  otherwise. The interaction graph is defined as  $G = (\mathcal{V}, \mathcal{E})$ , where  $\mathcal{V} = \mathcal{U} \cup \mathcal{I}$  is the set of nodes, and  $\mathcal{E} = \{(i, j) | \forall i \in \mathcal{U}, \forall j \in \mathcal{I}, m_{ij} = 1\}$  is the set of edges.

A graph-based encoder utilizes the interaction bipartite graph to enhance the ID embedding quality. Unlike a table encoder which independently encodes the users and items, when generating the user and item embeddings, a graph encoder additionally leverages their graph relationships (e.g., co-purchase, shared interests, etc). We depict the graph encoder in Figure 1 (b).

Most existing graph-enhanced CF methods in the retrieval stage focus on scenarios without the incorporation of contextual features, such as NGCF [40] and LightGCN [12]. However, whether or not their effectiveness can be transferred to existing ranking methods needs further investigation. Since the message passing mechanism [12, 17, 40] in graph-enhanced CF methods is the key to extracting graph knowledge, a natural way to extend ranking methods with graph knowledge is to include this mechanism in their paradigms. Following this path, we adapt a well-studied linear message passing mechanism [12] to existing ranking methods.

Specifically, let  $f_g(\cdot, \cdot) : \mathcal{G} \times x \rightarrow \mathbb{R}^d$  be the graph encoder. For user  $i$  and item  $j$ ,  $f_g(\cdot, \cdot)$  conducts message passing in each layer to propagate and aggregate information from the neighborhood. The graph-encoded embedding for node  $i$  (user or item) is formulated as follows:

$$\mathbf{z}_i = f_g(\mathcal{G}, x_i) = \sum_{l=0}^L a_l \mathbf{z}_i^{(l)}, \quad (7)$$

where  $\mathbf{z}_i^{(l)} = \sum_{v \in N_i} \frac{1}{\sqrt{|N_i|} \sqrt{|N_j|}} \mathbf{z}_j^{(l-1)}$  and  $\mathbf{z}_j^{(0)} = f(x_j)$ , (8)

In Equation (8),  $\mathbf{z}_i^{(l)}$  is the embedding for node  $i$  in layer  $l$ ,  $N_i$  is the set of neighbors for node  $i$  in  $\mathcal{G}$ , and  $a_l$  is the readout coefficient for each layer- $l$ 's embeddings. With the obtained graph-based user and item ID embeddings, we can construct the input features following Equation (2). These input features can further be fed into any ranking method (e.g., DCN) to predict the interaction probabilities/ranking scores.

**Table 1: Comparative ranking performance of DCNV2 [39] with a table encoder (Tab.), those of DCNV2 with a graph encoder (Graph), the relative change ( $\Delta\%$ ), and time in seconds per training epoch and inference.**

Metric	Recall@10 ↑			Recall@20 ↑		
Dataset	Tab.	Graph	%Δ	Tab.	Graph	%Δ
ML-1M	10.73	11.72	9.25	16.81	17.93	6.66
Yelp	4.25	4.30	1.18	9.87	10.23	3.59
AmzBook	3.28	3.73	13.79	7.53	8.30	10.23
Anime	15.73	16.97	7.92	23.76	24.91	4.84
Metric	NDCG@10 ↑			NDCG@20 ↑		
Dataset	Tab.	Graph	%Δ	Tab.	Graph	%Δ
ML-1M	10.87	12.31	13.26	12.59	13.93	10.59
Yelp	2.20	2.16	-1.91	3.75	3.79	1.23
AmzBook	1.92	2.09	8.98	3.14	3.43	9.23
Anime	13.90	15.73	0.13	16.41	18.07	10.09
Metric	Time / Train Epoch ↓			Inference Time (s) ↓		
Dataset	Tab.	Graph	%Δ	Tab.	Graph	%Δ
ML-1M	1.81	4.06	124.31	0.10	0.24	140.00
Yelp	3.53	15.18	330.03	0.17	0.88	417.65
AmzBook	5.56	29.53	431.12	0.25	1.76	604.00
Anime	9.05	58.30	544.20	0.47	3.52	648.94

### 2.3 The Benefit of Graphs to Ranking Methods

In comparison with user and item ID embeddings obtained from a table encoder as described in Section 2.1, those obtained from a graph encoder as introduced in Section 2.2 possess additional graph topological knowledge. To empirically verify the benefit of such graph knowledge to CTR methods, we design an experiment where all comparison models are identical except their encoders. Specifically, we compare the ranking performance of a DCNV2 [39] with a table encoder and that of a DCNV2 with a graph encoder. Since the graph encoder additionally incorporates graph knowledge into ID embeddings, and these two frameworks only differ in the encoding, the ranking performance gap can indicate how integrating graph knowledge affects the ranking method (i.e., DCNV2).

In this experiment, we train models on four benchmark datasets, including Yelp2018 [11], Amazon-Books [43], MovieLens-1M [7], and Anime [4]. We evaluate the models' ranking ability with two ranking-based metrics, recall and NDCG, where larger values indicate superior ranking ability. Their results are the averaged performance under five random seeds and are shown in Table 1, where the **Tab.** columns represent results of the DCNV2 with a table encoder, and **Graph** columns refer to those of the DCNV2 with a graph encoder. Compared with DCNV2 equipped with a table encoder, we observe that DCNV2 equipped with a graph-based encoder consistently surpasses its counterpart across the four datasets and all metrics. The enhancement suggests that the additional graph knowledge incorporated in the graph-based ID embeddings helps improve the ranking abilities of the target ranking method.

However, this integration incurs a noticeable increase in computational overhead. It incurs on average  $\sim 480\%$  more overheads for the total training time, and  $\sim 605\%$  more for the total testing time. These excess computational overheads arise from the message passing operations described in Equation (8) – to acquire the ID embedding of a user/item, the model is required to query representations of all nodes within the 2-hop neighborhood of the node to conduct the further aggregation in between. Moreover, this phenomenon can be further aggravated on dense and large graphs where the average number of neighbors per node is large. For example, it encounters  $\sim 1049\%$  more time in training and  $\sim 1261\%$  more in testing in the Anime dataset. Therefore, in industrial applications where billions of users and items construct a massive graph, simply substituting the table encoder with a graph-based encoder is prohibitively expensive and hence impractical.

### 2.4 A Simple yet Effective Solution: Non-parametric Graph Convolution

Although the introduced graph knowledge helps improve the ranking ability of the methods, significant computational overheads come along as well. The majority of computational overheads are brought by **training with the graph encoder**, where the encoder repetitively performs the computationally expensive message-passing operation on every iteration.

To address the acute problem of the growth of computational resources, we divert the integration of graph knowledge from the training phase to the test time, with the proposed strategy shown in Figure 1 (d). Injecting graph knowledge at testing time enjoys two benefits: (i) It obviates the forward passing and backpropagation entailed by message passing during training in Equation 8, whose computational overhead increases quadratically wrt the dataset density [10, 48]; (ii) It avoids the computational overheads brought by repetitively performing message passing during training since it only performs message passing once at testing time. Our strategy can be decoupled into the following four steps:


#### Step 1: Constructing Similarity Matrices

The first step is to construct two similarity matrices within users and items based on graph knowledge. We denote the similarity matrix within users as  $\mathbf{A}_u$ , and the matrix within items as  $\mathbf{A}_i$ , where both depict the co-purchase relationship. The two matrices are formulated based on the interaction matrix  $\mathbf{M}$ :

$$\begin{aligned} \mathbf{A}_u &= \mathbf{D}_u^{-\frac{1}{2}} \hat{\mathbf{A}}_u \mathbf{D}_u^{-\frac{1}{2}}, \hat{\mathbf{A}}_u = \mathbf{M} \mathbf{M}^\top, \\ \mathbf{A}_i &= \mathbf{D}_i^{-\frac{1}{2}} \hat{\mathbf{A}}_i \mathbf{D}_i^{-\frac{1}{2}}, \hat{\mathbf{A}}_i = \mathbf{M}^\top \mathbf{M}, \end{aligned} \quad (9)$$

where  $\mathbf{D}_u \in \mathbb{R}^{|\mathcal{U}| \times |\mathcal{U}|}$  and  $\mathbf{D}_i \in \mathbb{R}^{|\mathcal{I}| \times |\mathcal{I}|}$  are diagonal matrices with  $\mathbf{D}_{u/i}[k, k] = \sum_j \mathbf{A}_{u/i}[k, j]$ . Intuitively,  $\hat{\mathbf{A}}_u[i, j]$  represents the number of items interacted with both user  $i$  and  $j$ . Similarly,  $\hat{\mathbf{A}}_i[k, j]$  represents the number of users interacted with both item  $k$  and  $j$ .

For popular users/items associated with massive interactions, their corresponding entries in  $\hat{\mathbf{A}}_{u/i}$  tend to numerically dominate the corresponding entries of unpopular users/items associated with relatively fewer interactions. In other words, the similarity scores between popular users/items are consistently larger than those between unpopular users/items.



**Figure 2: An example of calculating the weights for aggregating the inference results of the constructed user-item pairs. The top-2 similar users and items are retrieved to construct  $2 * 2 = 4$  user-item pairs.**

To remove the bias caused by node popularity (i.e., node degree), we further normalize  $\hat{\mathbf{A}}_{u/i}$  by accounting for the varying degree of each user/item, and denote the normalized similarity matrices as  $\mathbf{A}_{u/i}$ . This normalization step prevents nodes (users or items) with disproportionately high degrees from occupying overly high similarity scores [17]. Therefore, we regard  $\mathbf{A}_{u/i}$  as the similarity matrices of users/items.

### Step 2: Retrieving Relevant User/Item Candidates

Relevant users and items based on the similarity matrices  $\mathbf{A}_{u/i}$  are retrieved in the second step. Specifically, for user  $i$ , we obtain  $n_k$  users with the corresponding top- $n_k$  similarity scores in  $\mathbf{A}_u[i]$ , denoted as  $\mathcal{U}_i$ , where  $\mathbf{A}_u[i]$  refers to the  $i$ -th row of  $\mathbf{A}_u$ . Similarly, for item  $j$ , we obtain  $n_k$  items with the corresponding top- $n_k$  similarity scores in  $\mathbf{A}_i[j]$ , denoted as  $\mathcal{I}_j$ . The corresponding  $2n_k$  similarity scores in  $\mathbf{A}_{u/i}$  are extracted to further calculate the aggregation weights of the later constructed user-item pairs. Intuitively, the higher the user/item similarity score is, the more aggregation weights should be assigned to the involved user-item pairs.

### Step 3: Constructing Relevant User-item Pairs

Since the users and items are selected based on the similarities in  $\mathbf{A}_{u/i}$ , which is constructed based on the collaborative filtering signals in  $\mathbf{M}$ , the interaction signals between the selected users and items naturally contain structural knowledge. Consequently, a set of relevant user-item pairs is constructed by combining each user in  $\mathcal{U}_i$  with each item in  $\mathcal{I}_j$ . This set contains  $n_k^2$  pairs and is defined as  $\{\mathcal{M}_{ij} : (u, v) | \forall u \in \mathcal{U}_i, \forall v \in \mathcal{I}_j\}$ . For each user-item pair in  $\mathcal{M}_{ij}$ , we calculate its weight as the multiplication of the corresponding user and item similarity scores, as shown in the middle of Figure 2. We denote the weight matrix for all user-item pairs in  $\mathcal{M}_{ij}$  as  $\text{Weight}(\mathcal{M}_{ij})$ , where  $\text{Weight}(\mathcal{M}_{ij})[u, v]$  represents the weight for the user  $u$  and item  $v$  pair.

### Step 4: Label Aggregation Among Relevant Ranking Scores

After obtaining the  $n_k^2$  user-item pairs, the target ranking model is queried for the corresponding  $n_k^2$  inference results, which are

aggregated based on each pair's weight  $\text{Weight}(\mathcal{M}_{ij})[u, v]$ :

$$p'_{ij} = \frac{\sum_{(u,v) \in \mathcal{M}_{ij}} \text{Weight}(\mathcal{M}_{ij})[u, v] * r(\mathbf{z}_{uv})}{\sum_{(u,v) \in \mathcal{M}_{ij}} \text{Weight}(\mathcal{M}_{ij})[u, v]}, \quad (10)$$

$$\mathbf{z}_{uv} = [f(x_u) \| f(x_v) \| h(\mathbf{c}_{uv})], \quad (11)$$

where  $r(\cdot)$  refers to an arbitrary trained ranking model such as the one we describe in Section 2.1. The aggregated inference result  $p'_{ij}$  is the final re-ranking score.

Although  $\text{Weight}(\mathcal{M}_{ij})$  is feasible to aggregate the inference results, we notice that the contribution proportion of the most similar user-item pair (i.e., pairs constructed with the most similar user and item) is too small. For the example in Figure 2, the most similar pair only contributes to  $1/(1 + 0.25 + 0.5 + 0.125) \approx 53\%$  of the final result.

This overly small proportion may deviate the final result too much from the prediction of the original user-item pair, and thereby downgrade the performance. To resolve this issue, we adopt a pairwise aggregation mechanism: (i) *First*, normalize the weights in  $\text{Weight}(\mathcal{M}_{ij})$  by dividing them by the maximum value in the matrix. This step normalizes the values in  $\text{Weight}(\mathcal{M}_{ij})$  between 0 to 1; (ii) *Second*, modify the aggregation weight for the most similar pair (i.e., the pair with the weight value as 1) by considering aggregating the pairwise inference result between this pair and the other pairs, as follows:

$$\text{Weight}(\mathcal{M}_{ij})[u', v'] = \sum_{(u,v) \neq (u',v')} (1 - \text{Weight}(\mathcal{M}_{ij})[u, v]), \quad (12)$$

$$\text{where } u', v' = \arg \max_{(u,v)} \text{Weight}(\mathcal{M}_{ij})[u, v].$$

The final re-ranking score is derived using Equation 11 with the maximum entry in  $\text{Weight}(\mathcal{M}_{ij})$  modified as above. In the example in Figure 2, the proportion of the most similar pair is modified to  $(1 - 0.25) + (1 - 0.5) + (1 - 0.125) = 2.125$ . This modified maximum weight is approximately 71% of the sum of the weights of all pairs, which is higher than 53% before its modification. In the following experiment section, we validate this design and empirically demonstrate how the proportion changes as  $n_k$  increases.

## 3 Experiment

### 3.1 Setup

**3.1.1 Datasets.** We select four publicly available recommendation benchmark datasets for the experiments, including Yelp2018 [11] and Amazon-Books [43] (relatively sparse), as well as MovieLens-1M [7] and Anime [4] (relatively dense). The statistics of the datasets are shown in Table 3. For all datasets, we convert explicit user-to-item ratings to binary labels through thresholding. We randomly split datasets with a ratio of 0.8/0.1/0.1 for training, validation, and testing, respectively.

**3.1.2 Baselines.** We select seven models as our baselines: NFM [11], DeepFM [8], xDeepFM [21], DCN [38], DCCN [39], AutoInt [32], and EulerNet [34]. NFM [11], DeepFM [8], and xDeepFM [21] combine the advantages of Factorization Machines [29] and deep neural networks to capture complex non-linear and high-order feature interactions. DCN [38, 39] learns explicit and implicit features

**Table 2: Ranking performance comparison by applying our approach to the baselines. ‘Original’ and ‘+Ours’ represent performance of the original models and those applied with our strategy respectively. ‘%Δ’ denotes the relative performance change.**

Model	NDCG@10			NDCG@20			Recall@10			Recall@20		
	Original	+Ours	%Δ	Original	+Ours	%Δ	Original	+Ours	%Δ	Original	+Ours	%Δ
ML-1M												
NFM	9.42	9.44	0.21	10.88	10.96	0.74	8.97	9.03	0.69	14.26	14.50	1.67
DeepFM	10.17	10.41	2.32	11.70	12.02	2.72	9.61	9.98	3.87	15.08	15.68	3.98
xDeepFM	8.08	8.17	1.06	9.65	9.82	1.72	8.04	8.19	1.92	13.13	13.52	3.00
DCN	10.07	11.19	11.08	11.58	12.73	9.98	9.47	10.17	7.33	14.97	16.11	7.64
DCNV2	10.76	10.91	1.39	12.26	12.52	2.10	10.02	10.26	2.39	15.55	16.12	3.69
AutoInt	8.72	9.18	5.23	10.32	10.82	4.84	8.66	9.19	6.07	13.96	14.60	4.60
EulerNet	8.91	9.11	2.34	10.51	10.96	4.32	8.73	9.24	5.86	14.08	14.83	5.28
Yelp2018												
NFM	2.92	3.35	14.58	4.58	5.10	11.35	5.28	6.06	14.77	11.23	12.30	9.58
DeepFM	1.95	2.18	11.78	3.52	3.83	8.63	3.90	4.33	10.86	9.66	10.32	6.86
xDeepFM	2.53	2.76	9.34	4.12	4.42	7.28	4.68	5.11	9.19	10.45	11.09	6.11
DCN	2.20	2.46	11.73	3.75	4.12	10.04	4.25	4.75	11.78	9.87	10.75	8.93
DCNV2	2.07	2.30	10.81	3.63	3.94	8.42	4.06	4.49	10.55	9.77	10.47	7.16
AutoInt	1.98	2.20	11.20	3.53	3.82	8.33	3.92	4.32	10.37	9.57	10.21	6.64
EulerNet	3.39	3.68	8.74	5.08	5.44	7.09	5.84	6.34	8.59	11.90	12.61	5.93
Amazon-Books												
NFM	2.67	3.04	13.69	4.06	4.47	10.25	4.52	5.09	12.42	9.10	9.81	7.85
DeepFM	2.40	2.65	10.69	3.69	4.01	8.50	4.00	4.43	10.86	8.41	9.01	7.04
xDeepFM	2.24	2.46	9.74	3.46	3.75	8.32	3.69	4.08	10.62	7.88	8.46	7.34
DCN	1.91	2.20	15.15	3.13	3.52	12.40	3.28	3.81	16.23	7.53	8.27	9.94
DCNV2	2.42	2.68	10.65	3.76	4.09	8.72	4.16	4.60	10.52	8.67	9.30	7.31
AutoInt	2.35	2.57	9.63	3.65	3.94	7.89	3.93	4.30	9.47	8.36	8.92	6.75
EulerNet	2.46	2.64	7.57	3.68	3.91	6.37	3.93	4.30	9.30	8.08	8.57	6.14
Anime												
NFM	14.74	14.68	-0.37	17.60	17.69	0.55	17.33	17.71	2.19	26.11	26.92	3.10
DeepFM	14.22	14.58	2.56	16.55	17.59	6.30	17.09	17.70	3.61	26.02	26.85	3.20
xDeepFM	15.09	15.08	-0.01	18.19	18.30	0.60	18.12	18.36	1.34	27.50	28.11	2.23
DCN	13.90	14.17	1.94	16.41	17.03	3.77	15.73	16.50	4.90	23.76	25.26	6.32
DCNV2	14.98	15.30	2.10	17.66	18.12	2.62	17.08	17.73	3.79	25.65	26.72	4.19
AutoInt	12.32	13.27	7.66	15.00	16.05	7.04	14.59	15.99	9.54	22.75	24.44	7.45
EulerNet	12.91	13.42	3.95	15.39	16.03	4.20	14.86	15.73	5.83	22.74	23.96	5.37

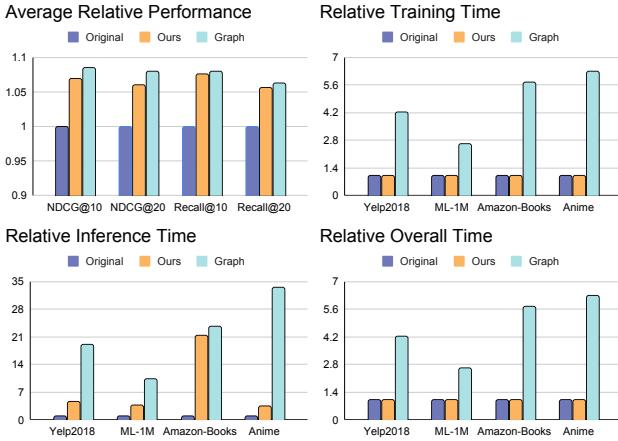
through a cross-network and a deep neural network, respectively. AutoInt [32] utilizes self-attentive neural networks to learn more effective feature interactions. EulerNet [34] learns high-order interaction features by transforming their exponential powers into linear combinations of the modulus and phase of complex features.

**3.1.3 Training.** We employ the AdamW optimizer for optimization and adopt binary cross-entropy as the loss function to train the models on the training set. We run a fixed number of grid searches over all the baseline models’ provided hyper-parameters for their best AUC performance on the validation set for the CTR prediction task. With the best hyper-parameters, we train the models under five random seeds and save all the checkpoints. The training and inference processes are conducted on an NVIDIA RTX 3090 GPU

**Table 3: The statistics of four benchmark datasets.**

Dataset	#User	#Item	#Interaction	Sparsity
ML-1M	6,041	3,261	998,539	0.9493
Yelp	77,278	45,639	2,103,896	0.9994
AmzBook	68,498	65,549	2,954,716	0.9993
Anime	55,119	7,364	6,270,078	0.9846

with 24 GB of memory, and the user-user and item-item similarity matrices are pre-computed on a standard commercial CPU with 128 GB of RAM. We adopt Recbole [49] to conduct all the experiments.




**Figure 3:** The relative performance is averaged over the four benchmark datasets and is scaled based on the original model performance. The overall time is the summation of training and inference time. **Original** denotes the original method, **Ours** denotes methods applied with our strategy, and **Graph** denotes the naive framework mentioned in Sec. 2.2.

**3.1.4 Evaluation.** For each baseline model, we evaluate our strategy with the inference results from the previously trained and saved model checkpoints. We tune the number of candidates  $n_k$  with the same amount of grid-searches and re-evaluate the model performance. The ranking ability is evaluated by two ranking-based metrics, NDCG@K and Recall@K, both of which assign higher scores to models with stronger ranking abilities. All reported results are averaged over the results under the five random seeds.

### 3.2 Ranking Performance Improvement

We evaluate the baseline models’ ranking performance on the four benchmark datasets, and denote the results under the **Original** columns in Table 2. We apply our strategy to each of the baseline models, and report the re-evaluated ranking performance under the **+Ours** columns in Table 2. The relative performance change under the **%Δ** columns is wrt the original method.

From the table, we observe that: (i) our strategy stably improves the performance over the original model, and only demonstrates a slight performance downgrade in some rare cases (i.e., NFM and xDeepFM in Anime). These results validate that our strategy is generally effective in improving ranking performance across various ranking models and benchmark datasets. (ii) The relative improvements are relatively evident in sparse datasets (i.e., Yelp2018, Amazon-Books) than dense datasets (i.e., ML-1M, Anime). This is because, in sparse datasets where the number of interactions associated with each node is relatively small, the node embeddings receive fewer collaborative filtering signals from their neighbors during training. Therefore, at test time, our strategy can compensate for insufficient training of the user/item embeddings with the injected graph knowledge. In contrast, user/item embeddings trained in dense datasets receive more training signals from a larger number



**Figure 4:** The ranking performance of DCN to varied number of considered neighbors  $n_k$  in our strategy.

of neighbors, resulting in less space for improvement at testing time by strategy.

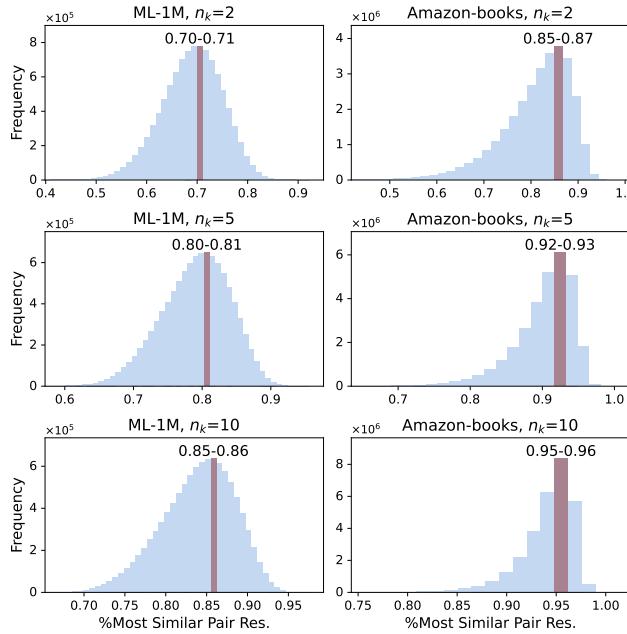
### 3.3 Time Efficiency

To demonstrate the efficiency of our strategy, we again take DCN as a typical ranking model to compare the averaged training, inference, and overall time in the three settings: DCN, DCN enhanced with our strategy, and the naive graph-enhanced DCN. The results shown in Figure 3 indicate that applying our strategy to DCN achieves a performance comparable to that of the naive graph-enhanced DCN, with only less than 2% of extra time overall. Applying our strategy to a well-trained CTR method does not introduce additional training time, and only quadratically increases the inference time wrt  $n_k$ . For each dataset, the most relevant  $n_k$  users and items neighbors can be precomputed, making the corresponding computational overhead one-off relative to a dataset.

### 3.4 Effect of $n_k$ to our strategy

To analyze how the number of relative user/item candidates  $n_k$  in our strategy affects the performance, we apply our strategy to DCN with varied  $n_k$  in  $\{1, 2, 5, 10\}$  to compare the performance. The results are shown in Figure 4. From the figure, we see that our strategy, when applied to dense graphs such as ML-1M and Anime, demonstrates improved performance with smaller  $n_k$ ’s. In contrast, on sparse graphs such as Yelp2018 and Amazon-books, our strategy yields better results with larger  $n_k$ ’s. This is because the value of  $n_k$  controls the range of the neighborhood considered for graph knowledge extraction. When  $n_k$  is small, the extracted graph knowledge is sufficient to improve CTR methods trained on dense graphs but insufficient for those trained on sparse graphs.

We further analyze how  $n_k$  affects the resultant contribution proportion of the most similar user-item pair in our strategy. Specifically, we apply our strategy to DCN with varied  $n_k$  in  $\{2, 5, 10\}$ ,



**Figure 5: Contribution proportion distribution for the most similar user-item pair, where most common proportion values are highlighted and denoted above.**

and depict the contribution proportion distribution of the most similar user-item pair in Figure 5. From the figure, we see that as  $n_k$  increases (i) the contribution proportion for the most similar user-item pair increases, and (ii) the distribution is more concentrated (i.e., it spans fewer values). Intuitively, the two phenomena suggest that within a close neighborhood (i.e.,  $n_k$  is small), our strategy adjusts the contribution proportion of the most similar pair in a wider range with relatively smaller numerical values. Conversely, within an extensive neighborhood (i.e.,  $n_k$  is large), our strategy conservatively adjusts the proportions in a tighter range with relatively larger values.

## 4 Related Work

A real-world RecSys consists of roughly two stages: candidate retrieval and item ranking. In the retrieval stage, *collaborative filtering* (CF) is commonly utilized, whereas in the ranking stage where rich contextual information is additionally incorporated, *click-through rate* (CTR) models are widely adopted in this stage.

**Collaborative Filtering.** As a prevalent technique that is widely employed in modern RecSys, CF makes recommendations based on the idea that similar users tend to have similar preferences [41]. Traditional CF methods aim to reconstruct user-item interactions with parameterized user and item embeddings. They model the reconstruction as a matrix factorization process with the user and item ID embeddings [18, 30, 47]. Some other methods maintain the ID embeddings and adopt neural networks to enhance the interaction modeling in between [11, 33]. Apart from improving interaction modeling, other recent works focus on refining other aspects, such

as the objectives and learning paradigms for performance enhancement [3, 20, 37, 46].

**Graph-based Collaborative Filtering.** Apart from signals from direct interactions, high-order CF signals in the user-item bipartite graph are crucial for personalized recommendations as well. These signals can be captured by the graph convolution operation in most graph neural networks (GNNs) [9, 17, 36]. Prior efforts adopt GCN [17] to the user-item interaction graph [1, 40, 44] to capture CF signals in the neighborhood. Later on, LightGCN [12] simplifies the graph convolution in GCN by preserving only linear neighborhood aggregation.

In addition to improving the structure of GNNs, recent works [2, 22, 24, 43, 45] enforce contrastive learning constraints in training the models for improved performance. For example, SGL [43] performs classical graph augmentation to the original bipartite graph to reinforce node representation learning via self-discrimination. SimGCL [45] refines the graph augmentation strategy in SGL with the perturbation of uniform noises and contrasts between the two perturbed graph views. NCL [22] explicitly incorporates potential neighbors into constructing contrastive pairs and defines a structure-contrastive objective to optimize.

**Click-through Rate Prediction.** As a widely adopted task for training models in the ranking stage, CTR prediction is defined as predicting the interaction likelihood between a user and an item given the user ID, item ID, and optional context features as input. The incorporated contextual information includes user demographic features, item description, interaction timestamps, etc. These features additionally consider the variability of user preferences for items across different contexts in which they interact with the system.

Early works in CTR seek efficient interactions between the interaction and the contextual information. They preserve low-order feature interactions through the prominent Factorization Machines [29] and seek high-order feature interactions through deep neural networks (DNNs) [8, 11, 21]. Later works modify the layer design within a deep neural network to automatically learn bounded-degree feature interactions [38, 39].

Some recent studies project the features to other predefined hyperspaces [32, 34] for more efficient and complex feature interactions, and the other [28] improves the model robustness from the supervision perspective.

Our strategy follows the established paradigm of CTR ranking methods. Unlike traditional graph-based CF approaches that integrate graph structure during training, our method leverages graph knowledge exclusively at inference time. This utilization of graph connectivity avoids the overhead of graph-based training while still benefiting from structural signals, making it lightweight and easily integrable into existing pipelines.

## 5 Conclusion

In this work, we investigate how to efficiently leverage graphs to improve the performance of models in the ranking stage of a RecSys, where rich contextual information is utilized.

We first demonstrate a naive graph-enhanced ranking framework, where graph knowledge is incorporated in the encoder of a ranking method via the message-passing operation. While this framework is empirically effective in improving ranking performance, the substantial computational overheads entailed by training with a graph encoder render this framework prohibitively expensive for real-world applications.

In light of this, we propose a non-parametric graph convolution strategy for ranking methods that utilizes graphs only once at test time to improve their ranking abilities. Our strategy can be used as a plug-and-play module, and can be easily employed with various ranking methods with little to no additional computational cost.

We conduct comprehensive experiments across four benchmark datasets with various densities to demonstrate that our strategy brings noticeable ranking performance improvements (i.e., 8.1% on average) during testing time with little to no additional computational overheads (i.e., 0.5% on average).

## Acknowledgments

This work was partially supported by the NSF under grants IIS-2321504, IIS-2334193, IIS-2217239, CNS-2426514, and CMMI-2146076, ND-IBM Tech Ethics Lab Program, Dartmouth College, and Notre Dame Strategic Framework and Poverty Initiative Research Grants (2025). Any expressed opinions, findings, conclusions, or recommendations are those of the authors and do not necessarily reflect the views of the sponsors.

## References

- [1] Rianne van den Berg, Thomas N Kipf, and Max Welling. 2018. Graph convolutional matrix completion. In *ACM SIGKDD Conference*.
- [2] Xuheng Cai, Chao Huang, Lianghao Xia, and Xubin Ren. 2022. LightGCL: Simple Yet Effective Graph Contrastive Learning for Recommendation. In *International Conference on Learning Representations*.
- [3] Chong Chen, Min Zhang, Yongfeng Zhang, Yiqun Liu, and Shaoping Ma. 2020. Efficient neural matrix factorization without sampling for recommendation. *ACM Transactions on Information Systems* (2020).
- [4] Cooper Union. 2023. Anime Recommendations Database. <https://www.kaggle.com/datasets/CooperUnion/anime-recommendations-database>. Accessed on: 2023-11-11.
- [5] Wensi Fan, Yao Ma, Qing Li, Yuan He, Eric Zhao, Jiliang Tang, and Dawei Yin. 2019. Graph neural networks for social recommendation. In *ACM Web Conference*.
- [6] Carlos A Gomez-Uribe and Neil Hunt. 2015. The netflix recommender system: Algorithms, business value, and innovation. *ACM Transactions on Management Information Systems* (2015).
- [7] GroupLens Research. 2023. MovieLens 1M Dataset. [https://grouplens.org/datasets/movielens/1m/](https://grouplens.org/datasets/movielens/). Accessed: 2023-11-11.
- [8] Hufeng Guo, Ruiming Tang, Yuning Ye, Zhenguo Li, and Xiuqiang He. 2017. DeepFM: a factorization-machine based neural network for CTR prediction. In *International Joint Conference on Artificial Intelligence*.
- [9] Will Hamilton, Zhitao Ying, and Jure Leskovec. 2017. Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems*.
- [10] Xiaotian Han, Tong Zhao, Yozhen Liu, Xia Hu, and Neil Shah. 2023. MLPInit: Embarrassingly Simple GNN Training Acceleration with MLP Initialization. In *International Conference on Learning Representations*.
- [11] Xiangnan He and Tat-Seng Chua. 2017. Neural factorization machines for sparse predictive analytics. In *ACM SIGIR Conference*.
- [12] Xiangnan He, Kuan Deng, Xiang Wang, Yan Li, Yongdong Zhang, and Meng Wang. 2020. LightGCN: Simplifying and powering graph convolution network for recommendation. In *ACM SIGIR Conference*.
- [13] Mohsen Jamal and Martin Ester. 2010. A matrix factorization technique with trust propagation for recommendation in social networks. In *ACM Recommender Systems conference*.
- [14] Clark Mingxuan Ju, Leonardo Neves, Bhuvesh Kumar, Liam Collins, Tong Zhao, Yuwei Qiu, Qing Dou, Sohail Nizam, Sen Yang, and Neil Shah. 2025. Revisiting Self-attention for Cross-domain Sequential Recommendation. In *ACM SIGKDD Conference*.
- [15] Clark Mingxuan Ju, Leonardo Neves, Bhuvesh Kumar, Liam Collins, Tong Zhao, Yuwei Qiu, Qing Dou, Yang Zhou, Sohail Nizam, Rengim Ozturk, et al. 2025. Learning Universal User Representations Leveraging Cross-domain User Intent at Snapchat. In *ACM SIGIR Conference*.
- [16] Mingxuan Ju, William Shiao, Zhichun Guo, Yanfang Ye, Yozhen Liu, Neil Shah, and Tong Zhao. 2024. How Does Message Passing Improve Collaborative Filtering?. In *Advances in Neural Information Processing Systems*.
- [17] Thomas N Kipf and Max Welling. 2017. Semi-supervised classification with graph convolutional networks. In *International Conference on Learning Representations*.
- [18] Yehuda Koren, Robert Bell, and Chris Volinsky. 2009. Matrix factorization techniques for recommender systems. *IEEE Computer* (2009).
- [19] Yehuda Koren, Steffen Rendle, and Robert Bell. 2021. Advances in collaborative filtering. *Recommender systems handbook* (2021).
- [20] Dongha Lee, SeongKu Kang, Hyunjun Ju, Chanyoung Park, and Hwanjo Yu. 2021. Bootstrapping user and item representations for one-class collaborative filtering. In *ACM SIGIR Conference*.
- [21] Jianxun Lian, Xiaohuan Zhou, Fuzheng Zhang, Zhongxia Chen, Xing Xie, and Guangzhong Sun. 2018. xdeepfm: Combining explicit and implicit feature interactions for recommender systems. In *ACM SIGKDD Conference*.
- [22] Zihan Lin, Changxin Tian, Yupeng Hou, and Wayne Xin Zhao. 2022. Improving graph collaborative filtering with neighborhood-enriched contrastive learning. In *ACM Web Conference*.
- [23] Hao Ma, Haixuan Yang, Michael R Lyu, and Irwin King. 2008. Sorec: social recommendation using probabilistic matrix factorization. In *ACM International Conference on Information and Knowledge Management*.
- [24] Yunshan Ma, Yingzhi He, An Zhang, Xiang Wang, and Tat-Seng Chua. 2022. CrossCBR: cross-view contrastive learning for bundle recommendation. In *ACM SIGKDD Conference*.
- [25] Zhongyu Ouyang, Qianlong Wen, Chunhui Zhang, Yanfang Ye, and Soroush Vosoughi. 2025. Towards Human-like Preference Profiling in Sequential Recommendation. arXiv:2506.02261 [cs.IR]. <https://arxiv.org/abs/2506.02261>
- [26] Zhongyu Ouyang, Chunhui Zhang, Shifu Hou, Shang Ma, Chaoran Chen, Toby Li, Xusheng Xiao, Chuxu Zhang, and Yanfang Ye. 2024. Symbolic Prompt Tuning Completes the App Promotion Graph. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*.
- [27] Zhongyu Ouyang, Chunhui Zhang, Shifu Hou, Chuxu Zhang, and Yanfang Ye. 2024. How to improve representation alignment and uniformity in graph-based collaborative filtering?. In *Proceedings of the International AAAI Conference on Web and Social Media*.
- [28] Zhongyu Ouyang, Chunhui Zhang, Yaning Jia, and Soroush Vosoughi. 2025. Scaled supervision is an implicit lipschitz regularizer. In *Proceedings of the International AAAI Conference on Web and Social Media*.
- [29] Steffen Rendle. 2010. Factorization machines. In *IEEE ICDM*.
- [30] Steffen Rendle, Christoph Freudenthaler, Zeno Gantner, and Lars Schmidt-Thieme. 2009. BPR: Bayesian personalized ranking from implicit feedback. In *The Conference on Uncertainty in Artificial Intelligence*.
- [31] J Ben Schafer, Joseph Konstan, and John Riedl. 1999. Recommender systems in e-commerce. In *ACM conference on Electronic commerce*.
- [32] Weiping Song, Chence Shi, Zhiping Xiao, Zhijian Duan, Yewen Xu, Ming Zhang, and Jian Tang. 2019. AutoInt: Automatic feature interaction learning via self-attentive neural networks. In *ACM International Conference on Information and Knowledge Management*.
- [33] Yi Tay, Luu Anh Tuan, and Siu Cheung Hui. 2018. Latent relational metric learning via memory-based attention for collaborative ranking. In *ACM Web Conference*.
- [34] Zhen Tian, Ting Bai, Wayne Xin Zhao, Ji-Rong Wen, and Zhao Cao. 2023. EulerNet: Adaptive Feature Interaction Learning via Euler's Formula for CTR Prediction. In *ACM SIGIR Conference*.
- [35] Aaron Van den Oord, Sander Dieleman, and Benjamin Schrauwen. 2013. Deep content-based music recommendation. In *Advances in Neural Information Processing Systems*.
- [36] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Lio, and Yoshua Bengio. 2017. Graph attention networks. In *International Conference on Learning Representations*.
- [37] Chenyang Wang, Yuanqing Yu, Weizhi Ma, Min Zhang, Chong Chen, Yiqun Liu, and Shaoping Ma. 2022. Towards Representation Alignment and Uniformity in Collaborative Filtering. In *ACM SIGKDD Conference*.
- [38] Ruoxi Wang, Bin Fu, Gang Fu, and Mingliang Wang. 2017. Deep & cross network for ad click predictions.
- [39] Ruoxi Wang, Rakesh Shivanna, Derek Cheng, Sagar Jain, Dong Lin, Lichan Hong, and Ed Chi. 2021. Dcn v2: Improved deep & cross network and practical lessons for web-scale learning to rank systems. In *ACM Web Conference*.
- [40] Xiang Wang, Xiangnan He, Meng Wang, Fuli Feng, and Tat-Seng Chua. 2019. Neural graph collaborative filtering. In *ACM SIGIR Conference*.
- [41] Yinwei Wei, Wenqi Liu, Fan Liu, Xiang Wang, Liqiang Nie, and Tat-Seng Chua. 2023. Lightgt: A light graph transformer for multimedia recommendation. In *ACM SIGIR Conference*.

- [42] Jiancan Wu, Xiangnan He, Xiang Wang, Qifan Wang, Weijian Chen, Jianxun Lian, and Xing Xie. 2022. Graph convolution machine for context-aware recommender system. *Frontiers of Computer Science* (2022).
- [43] Jiancan Wu, Xiang Wang, Fuli Feng, Xiangnan He, Liang Chen, Jianxun Lian, and Xing Xie. 2021. Self-supervised graph learning for recommendation. In *ACM SIGIR Conference*.
- [44] Rex Ying, Ruining He, Kaifeng Chen, Pong Eksombatchai, William L Hamilton, and Jure Leskovec. 2018. Graph convolutional neural networks for web-scale recommender systems. In *ACM SIGKDD Conference*.
- [45] Junliang Yu, Hongzhi Yin, Xin Xia, Tong Chen, Lizhen Cui, and Quoc Viet Hung Nguyen. 2022. Are graph augmentations necessary? simple graph contrastive learning for recommendation. In *ACM SIGIR Conference*.
- [46] An Zhang, Leheng Sheng, Zhibo Cai, Xiang Wang, and Tat-Seng Chua. 2023. Empowering Collaborative Filtering with Principled Adversarial Contrastive Loss. In *Advances in Neural Information Processing Systems*.
- [47] Honglei Zhang, Fangyuan Luo, Jun Wu, Xiangnan He, and Yidong Li. 2023. LightFR: Lightweight federated recommendation with privacy-preserving matrix factorization. *ACM Transactions on Information Systems* (2023).
- [48] Shichang Zhang, Yozen Liu, Yizhou Sun, and Neil Shah. 2022. Graph-less Neural Networks: Teaching Old MLPs New Tricks Via Distillation. In *International Conference on Learning Representations*.
- [49] Wayne Xin Zhao, Shanlei Mu, Yupeng Hou, Zihan Lin, Yushuo Chen, Xingyu Pan, Kaiyuan Li, Yujie Lu, Hui Wang, Changxin Tian, et al. 2021. Recbole: Towards a unified, comprehensive and efficient framework for recommendation algorithms. In *ACM International Conference on Information and Knowledge Management*.